

OligoTox-Thrombocytopenia: a curated, per-measurement dataset of oligonucleotide platelet toxicity

Narrative document — NIH/NCATS Oligonucleotide Toxicity (OligoTox) Open Data Challenge, Phase 2 (Data Generation)

Endpoint: thrombocytopenia / platelet toxicity · **Dataset:** 259 oligonucleotides × 1,959 graded measurements from 70 primary sources · **Licence:** CC-BY 4.0 · **Repository:** navierol/Oligos → thrombocytopenia/

1. Executive summary

OligoTox-Thrombocytopenia is an openly-releasable, per-measurement dataset of platelet-toxicity signals for therapeutic oligonucleotides. Thrombocytopenia is one of the toxicities of interest named in the Challenge announcement, and — unlike hepatotoxicity — no comparable public, structured, per-compound resource for it exists. One row is one *oligo* × *model/subject* × *delivery* × *dose* × *readout*, carrying both the **design predictors** (sequence, backbone, phosphorothioate count, sugar chemistry, conjugate) and a **graded outcome** on an ordinal 0–3 scale, together with per-row provenance and rights status.

1,959 MEASUREMENTS	259 OLIGONUCLEOTIDES	1,453 HUMAN ROWS	70 PRIMARY SOURCES	23 CROSS-SPECIES COMPOUNDS
-----------------------	-------------------------	---------------------	-----------------------	----------------------------------

The endpoint is bimodal, and the dataset is built to preserve that

Oligonucleotide-associated thrombocytopenia is not one phenomenon. The literature consistently describes (i) a **common, mild, dose- and plasma-concentration-dependent** platelet decline that plateaus and is reversible, and (ii) a **rare, severe, idiosyncratic, immune-mediated** thrombocytopenia with treatment-emergent antiplatelet antibodies and counts below $25 \times 10^9/L$. A dataset that averages these together destroys the distinction a predictive model most needs. Inotersen's FDA label documents both modes in one compound — the basis of its Boxed Warning — and is the anchor against which the grading rubric was calibrated. The FDA clinical review for that product states the same typology explicitly ("Rare" vs "Common"), which is independent regulatory confirmation of the framing rather than our inference.

Positive and negative controls

Role	Exemplars	Evidence
Positive — severe mode	inotersen; volanesorsen	Boxed Warning; platelets $< 25 \times 10^9/L$ in 3 % with antiplatelet IgG in 3/3; one fatal intracranial haemorrhage. Volanesorsen: dose-limiting, $< 100 \times 10^9/L$ in 47 % vs 0 % placebo.
Positive — mild mode	olezarsen; ISIS 104838; the (AC) _n series	Mean platelet decline of 6–10 % without bleeding; in-vitro activation with monotonic dependence on PS count.
Negative — neutral backbone	eteplirsen, golodirsen, viltolarsen, casimersen	PMO morpholino backbone; zero platelet content across their FDA labels; a pooled meta-analysis reports an explicit zero for the PMO subgroup against a 21.2 % overall rate.
Negative — targeted conjugate	inclisiran, givosiran, lumasiran, vutrisiran, patisiran	Explicit reported zeros, e.g. <i>Platelets</i> $\leq 75 \times 10^9/L$: 0 (0.0 %) in the inclisiran FDA review; an FDA-generated platelet shift table for vutrisiran/patisiran.
Negative — isosequential	ODN2395 phosphodiester; fesomersen	Same base sequence as an active comparator, differing only in backbone or conjugate.
Negative — mechanistic	GPVI-deficient donor platelets	No binding or activation up to 10 μM — a receptor-level negative that isolates the mechanism.

The inventory is enumerated in `data/controls_inventory.csv`: **31 control records across 26 compounds**, comprising 4 isosequential backbone pairs, 4 chemistry-variant comparators, 15 vehicle or concurrent-control arms, and 4 intended-pharmacology comparators. Each row records what the control controls for, its permitted use and its prohibited use.

A control is not a clinical negative, and this dataset has the first and not the second. The distinction decides what may be claimed. A *control* is a comparator measured alongside an exposure — an isosequential phosphodiester arm, a vehicle arm, a receptor-deficient donor. A *qualified clinical negative* would be a compound licensed to carry a "clinically negative for thrombocytopenia" label. The project's scientific governance records **zero** qualified clinical negatives: the clean-negative rule is *not yet available* and all twenty-one audited candidates are ineligible. So the controls above are reported as controls, and no compound in this dataset is labelled clinically safe.

The discipline that follows: a compound whose label never mentions platelets is *not* a measured zero. Every grade-0 row is backed by a source reporting that platelet monitoring was performed and found unremarkable. Absence of mention was never converted into evidence of absence, and the code that swept the labels states this in its own output. The neutral-backbone and targeted-conjugate rows above rest on *explicit reported zeros* for that reason, not on silence — though their trial-grain coverage is thin and is logged as an open gap in `data/recovery_ledger.csv`.

2. Main findings and conclusions

2.1 The ordering the phosphorothioate hypothesis predicts is present in the curated data

Platelet effects are held to track **phosphorothioate (PS) content and protein binding** rather than the hybridisation target, mechanistically via glycoprotein VI. Without any modelling, the assembled data orders in the direction that hypothesis predicts:

What this ordering is, and what it is not. It is a *description of curated labels*, not an independent reproduction of the structure–activity relationship. Three dependencies have to travel with it, and an earlier version of this document asserted the result without them:

- **The indicator is a curator-assigned label.** Every grade in these means was assigned by this curation effort against a documented rubric, not read from a source as a severity grade. All 1,959 graded rows are referred to the project scientist as **provisional**.
- **The laboratory rows' grades derive from a mapping the project prohibits.** 523 in-vitro and ex-vivo rows carry a grade derived from a clinical severity scale; the governing rules forbid mapping in-vitro bins to clinical grades absent clinical validation. Those rows are referred for re-scaling and the ordering must be re-read once they are re-scaled.
- **The thinnest cell carries 1 human row(s).** The neutral-backbone comparator — the load-bearing negative control for PS dependence — rests on that, and its trial-grain coverage is logged as an open gap.

So the defensible statement is: *the curated labels order in the predicted direction, under an indicator that is itself provisional*. It is not evidence that the relationship holds independently of the labelling, and no causal or predictive claim follows from it.

Backbone	human rows	compounds	mean grade	PS linkages	mean grade
PMO_neutral	1	1	0.00	0	0.42
PS_P0_mix	39	5	0.10	13–16	1.21
full_P0	35	4	0.46	17–19	1.24
full_PS	1,190	43	1.10	20+	1.44

These are HUMAN rows only. Earlier versions of this table pooled human and animal evidence, which averages across species, exposures and assay systems and produces a number no reviewer should rely on. Restricting to human evidence does not weaken the relationship — it sharpens it: the PS-count gradient runs 0.42 → 1.21 → 1.24 → 1.44. Animal evidence is retained in full in `data/measurements_animal.csv` and influences no human figure. Cells backed by few rows are shown with their row and compound counts so the reader can discount them; `PMO_neutral` in particular rests on 1 human row(s).

Modality orders in the same direction, under the same three dependencies, in human evidence alone: PMO 0.00 < GalNAc-siRNA 0.09 < siRNA 0.25 < splice-switching 0.90 < ASO gapmer 1.00 < aptamer 1.55. That a curated dataset recovers a relationship the field established experimentally is the primary evidence that it is fit to train on.

2.2 Five controlled comparisons — the dataset's most distinctive content

Curation deliberately sought **matched pairs and titration series**: cases where one design variable moves while the others are held fixed. These let a model learn a mechanism rather than a correlation.

#	Comparison	Variable isolated	Result
1	ODN2395 phosphodiester vs phosphorothioate	backbone only (0 → 21 PS); identical 22-mer	PAC-1 binding 1.65 → 18.94 MFI on human platelets
2	the (AC) _n series, n = 5...11	length, hence PS count, sequence held constant	monotonic 2.12 → 16.79 MFI; threshold between 9 and 11 linkages
3	ISIS 416858 vs fesomersen	19 → 13 PS + GalNAc; identical sequence and 5-10-5 design	dose-dependent monkey platelet decline vs none
4	volanesorsen vs olezarsen	GalNAc conjugate and the ~20× lower dose it permits	dose-limiting (47 % < 100 × 10 ⁹ /L) vs a 6–10 % mean decline, no major bleeding
5	ISIS 972163 vs ISIS 972116	one nucleoside in a wing (cEt → 2'-MOE)	556 vs 325 K/μL in the same rat table — a 41 % difference

Comparison 4 required recovering olezarsen's sequence, which is not published as a string, by deterministic parse of the IUPAC chemical name in its FDA label. The parse returns AGCTTCTTGTCAGCTTTAT — identical to volanesorsen's independently published sequence, exactly as expected for the GalNAc-conjugated APOC3 analogue, with the stated molecular formula P20S19 independently fixing full-PS chemistry at 20 residues. Two derivation paths reaching the same string is meaningful evidence against transcription artefact.

2.3 Human and animal evidence are separated, and 23 compounds bridge them

The announcement singles out datasets "based on in vitro human systems or able to **extrapolate data between in vitro human systems and animal data**". The dataset is **74 % human** (1,453 of 1,959 rows), and the axis is a first-class derived column, `subject_class`, computed from study type × species and independently re-derived by the QC gate so it cannot drift.

Human	rows	Animal	rows
clinical (outcome records)	1,002	in vivo	429
in vitro (platelets, PRP, whole blood)	426	in vitro	55
ex vivo	25	ex vivo	13
total	1,453 (78 compounds)	total	497 (174 compounds)

Those are outcome RECORDS, not trials, and the difference is large. Several rows are commonly dose-band cells of a single supplementary table of a single pooled analysis. Resolving the 1,002 human clinical records to study grain yields **85 evidence units**: 56 typed as a trial, 52 with a registry identifier verified against ClinicalTrials.gov or EudraCT, and 24 pooled analyses that are *not* trials.

No single headline trial count is published, because none is defensible. Of the units typed as a trial, only **39** carry any measurement row at all — the rest are trial-grain *anchors*, identified so that pooled data can be attributed to a real study rather than floating free. Of those, 23 have a platelet endpoint demonstrably assessed under a defined exposure and observation window, and 22 remain once units are excluded in which a platelet change is the intended pharmacology or the treated disease — an anti-von-Willebrand aptamer that raises platelets in type 2B von Willebrand disease, and a telomerase inhibitor dosed into thrombocythaemia. That last figure is still a *curator classification*, not a qualification.

Those 22 were then **independently re-audited** against four explicit tests: that the cited basis quotes evidence platelets were *measured* rather than merely reported absent; that dose and duration are both stated; that a platelet-specific at-risk denominator exists; and that the basis cites a retrievable locus. **19 of 22 pass all four.** So the ladder reads **56 typed as a trial → 39 with measurements → 19 passing this endpoint's audit**, and the last of those is the figure to quote. None of the three is a count of *qualified* trials: qualification against the project's sign-off gates is the scientist's determination, not this dataset's, and it has not been made. The per-unit results are in `data/toxicity_denominator_audit.csv` and the full ladder in `data/study_counts.csv`.

Overlap is declared, not assumed away. 21 edges in `data/study_nesting_ledger.csv` record a trial that appears here individually and also inside a pooled analysis recorded here. Each was established by arithmetic agreement on arm sizes against the pooling source's own table — not by compound-name similarity, which produced several false overlap claims during resolution. Participant totals are therefore **not summable**, and the dataset says so rather than printing a tempting grand total.

Containing human rows and animal rows does not by itself demonstrate extrapolation. Having the *same compound* characterised on both sides does, and **23 compounds** qualify — inotersen (199 human / 22 animal rows), volanesorsen (172 / 18), mipomersen (114 / 31) among them, each with a published sequence. `data/bridge_human_animal.csv` presents them side by side with their human and animal grades. That set, not the raw totals, is what a cross-species model can be trained and validated on.

2.4 Mechanism must be labelled, not assumed

Imetelstat causes severe thrombocytopenia, but by **on-target telomerase-inhibitor myelosuppression**, not PS-backbone platelet binding. Left unflagged, its 93 high-grade rows made up 74 % of the zero-PS bucket and pushed that bucket's mean grade *above every phosphorothioate bucket* — inverting the headline result. It is mechanism-flagged and excluded from structure–activity analysis while remaining in the dataset. Similarly, volanesorsen's anti-*drug* antibodies carried no altered safety profile and must not be conflated with inotersen's anti-*platelet* IgG, which tracked with severe events.

3. How the data were produced

This is an **in-silico curation** of pre-existing, publicly reported data — not wet-lab-generated data. "Experimental design" therefore describes source identification, extraction, harmonisation, grading, provenance and quality control. Full detail is in the accompanying methodology document.

3.1 Design

- **Grain:** one row per oligo × model/subject × delivery × dose/concentration × readout. A compound measured for platelet count *and* P-selectin at one dose yields two rows.
- **Two normalised tables** joined on `oligo_id`: `oligos.csv` (identity and design predictors, 20 columns) and `measurements.csv` (outcome and context, 24 columns). A denormalised analysis view is generated, never hand-edited.
- **Coverage goal:** all modalities, all study types, the full severity range, and **well-sourced negative controls** — which are what make the chemistry hypothesis falsifiable.
- Missing values are the literal string `TBD` — never guessed, never imputed as zero.

3.2 Acquisition

Sources were discovered by a **parallel multi-agent sweep across ten independent lanes** (clinical, regulatory, in-vitro human platelet, mechanism/immune, preclinical, non-ASO negatives, patents, reviews/meta-analyses, megakaryocyte models, sequence–chemistry panels), each blind to the others. Lanes that overlapped provided cross-confirmation; lanes that did not surfaced sources the others missed. Sources were ranked by **extractable per-measurement content**, not topicality, and each had to be *fetched* before being proposed. Primary sources were retrieved directly: NCBI E-utilities and Europe PMC full text plus supplementary files, the DailyMed SPL API for FDA labels, `accessdata.fda.gov` for FDA review documents, the EMA for SmPCs and EPAR assessment reports, and USPTO/Google Patents for patent full text and sequence listings.

Regulatory review documents proved the highest-value and least-mined vein: FDA multidiscipline reviews and EMA assessment reports carry per-patient platelet nadir listings, antibody cross-tabulations and dose-banded incidence grids that the corresponding journal publications do not.

3.3 Computational processing

The path from extraction to published CSV is fully scripted and re-runnable (`scripts/ingest_thrombo.sh`): harvest → merge lanes in a deterministic precedence order → assemble (apply verification verdicts, deduplicate on a natural key, assign stable keys) → re-apply committed verdicts → **QC gate** → rebuild the derived view and the generated documentation sections. QC gates the round: if it fails, the derived artefacts are not regenerated, so a broken dataset cannot ship with fresh-looking documentation. The **raw curation record** — every agent's returned rows, the verification verdicts, the source sweep — is committed alongside the data, so the published tables are reproducible from their inputs rather than merely re-checkable against their citations. Rebuilding from that record reproduces both CSVs byte-for-byte.

3.4 Accuracy control

- **Adversarial verification.** Extracted rows are passed to an independent agent instructed to *refute* them against the cited source. Two blocks totalling 659 rows have completed: **628 confirmed, 31 corrected, 0 rejected**. For one block every value was recomputed from the paper's raw per-replicate workbook and every mean, SD and n reproduced exactly.
- **Held-out known-answer test.** The inotersen label was parsed independently and its figures recorded *before* any extraction ran; the assembled dataset matched all six.
- **Cross-dataset sequence agreement.** Sequences were checked against a sister dataset's records, which were validated there against WHO INN nomenclature and patent sequence listings: **zero conflicts**.

Four defects found and corrected, worth reporting because three were invisible to schema validation. (i) A cross-source **name collision**: ODN 2395 denotes a phosphorothioate reagent everywhere except one paper that made a phosphodiester variant under the bare name — 39 rows describing an active PS compound had attached to a record asserting `ps_count = 0`, undercutting the central hypothesis. (ii) The **mechanism confound** of §2.4. (iii) A **header-misalignment** in three patent tables that silently returned red-cell counts where platelet counts were intended in 35 of 55 rows — real haematology numbers, of plausible magnitude, from the correct study. (iv) A controlled-vocabulary typo. All four produced structurally valid data; they were caught by asking whether the assembled data still reproduced a relationship the field already knows, which is why that check runs every round.

4. Indicators, predictors, and their distributions

4.1 The graded indicator

Each measurement carries an ordinal **thrombocytopenia_grade (0–3)**, CTCAE-aligned for clinical rows with explicit in-vitro analogues so a bench readout and a trial outcome land on a comparable scale: **0** no signal at tested exposure · **1** mild/reversible, $\geq 100 \times 10^9/L$ or activation only at supra-clinical concentration · **2** moderate, $50\text{--}99 \times 10^9/L$ or activation at clinically relevant concentration · **3** severe, $< 50 \times 10^9/L$, CTCAE grade 4, antibody-mediated, bleeding attributable to thrombocytopenia, or discontinuation.

Two conventions are applied uniformly, both stated because they are the easiest places to introduce systematic error. An **incidence row is graded by the severity of the event**, not by how small the incidence is. **Bleeding is graded on attribution**: grade 3 requires haemorrhage attributable to thrombocytopenia, so a trial reporting mild bleeding at or below the placebo rate is not scored severe.

Grade	rows	Readout category	rows	Study type	rows
0 — none	852	platelet count	1,083	clinical	1,002
1 — mild	499	clinical outcome	226	in vitro	485
2 — moderate	388	platelet activation	196	animal in vivo	434
3 — severe	220	binding / aggregation	236	ex vivo	38
		immunogenicity	105		
		megakaryocyte / other	113		

Readouts deliberately emphasise the **functional and immunological** axes over viability: 1,083 platelet-count rows, 196 activation, 236 binding/aggregation, and 105 immunogenicity rows capturing the antiplatelet-antibody mechanism that distinguishes the severe mode.

4.2 Predictor distribution among the 259 tested oligos

Modality	n	Backbone	n	Conjugate	n	Stage	n
ASO gapmer	190	full PS	193	none	228	research panel	135
other / tool	45	PS/PO mix	11	GalNAc	15	preclinical	53
PMO	6	full PO	6	lipid	5	phase 1	24
aptamer	7	PMO neutral	6	PEG / other	4	phase 2–3	16
splice-switching	5	mixed / NA	4	TBD	7	approved	17
GalNAc-siRNA / siRNA	6	TBD	39			class review	8

Sequences are present for 200 of 259 oligos (77 %) and were never guessed: the remainder are proprietary, class-level aggregates, or development-code compounds with no published sequence. **67 distinct target genes** are represented. Provenance is complete — every row carries a source identifier, an exact locus (table, figure, claim or label section) and a licence class resolved **per publisher and per regulator**: 291 rows from US federal agency work, 241 from granted US patent text, 216 from EMA documents (reusable with attribution, *not* US public domain), 229 under CC-BY or CC0, 747 under a NonCommercial clause, 11 under NoDerivatives, and **224 on a proposed hold** where no open licence could be located. That last figure is a decision requested of the data owner, not a withdrawal: **zero rows withdrawn, zero rows cleared** by this curation effort. The ledger is a **project classification, not legal clearance**, and it addresses reuse of extracted facts rather than republication of source files.

5. The gap this addresses

Phase 1 of this Challenge was dominated by hepatotoxicity; over half the winning teams targeted liver models. Thrombocytopenia is a named toxicity of interest that **no public, structured, per-compound resource covers**, despite being the toxicity that carries a Boxed Warning on a marketed oligonucleotide, that was dose-limiting for another, and that has caused a trial death.

The gap is not that the information is secret — it is that it is **scattered and unusable**. The evidence sits in FDA multidiscipline reviews, EMA assessment reports, patent example tables and the supplementary files of individual papers, in incompatible formats, without harmonised severity grading, and without any linkage between a compound's chemistry and its platelet outcome. Nothing lets a modeller ask "what does phosphorothioate count do to platelets across every compound ever reported?"

This dataset closes that specific gap. It contributes:

- **Harmonised grading** across in-vitro, ex-vivo, animal and clinical evidence on one 0–3 scale, so sources that never shared a metric become comparable.
- **Design predictors joined to outcomes** — sequence, PS count, backbone, conjugate — which is the join no existing source performs.
- **A deliberately populated negative class** (795 grade-0 rows). Negative results are systematically under-published; sourcing them from regulatory documents, where monitoring is reported whether or not it was eventful, is a contribution in itself.
- **Controlled comparisons and a 23-compound cross-species bridge**, which are what make mechanism learnable rather than merely correlated.
- **Per-row rights tracking**, recording what each publisher or regulator declares, with the basis, so a consumer can filter to the rows matching their own reuse determination. It is a project classification, not legal clearance, and it is about reuse of extracted facts rather than republication of source files.

6. Using the data to build a predictive model

A model this dataset previously demonstrated has been withdrawn. Earlier versions of this narrative reported grouped ROC-AUC for a classifier predicting platelet effect (grade ≥ 1) from design features across 228 compounds. The project's scientific governance package classifies that exact model — a sequence-only clinical thrombocytopenia classifier — as **BLOCKED**, with training prohibited and no permitted claim. It is no longer run and its scores are retracted (`data/model_demo_results.json` now holds the retraction record). Five independently checkable reasons: it used 228 compounds where 16 are authorised for mechanistic work and 7 carry any *provisional* clinical eligibility; it grouped folds by compound rather than by exact sequence, so isosequential pairs such as volanesorsen/olezarsen and inotersen/eplotersen fell on both sides of a split; it used grade 0 rows as negatives although the qualified clinical-negative count is zero; it treated pooled dose-band rows as independent observations; and `is_human` ranked second of eighteen features, meaning the model was substantially learning which rows came from human studies.

What the dataset supports instead is reported in `data/approved_analyses.json`, generated by `scripts/model_demo.py`, restricted to the lanes the governance package authorises:

Lane	Status	Result
Matched-sequence mechanistic contrasts	approved, descriptive	6 contrasts; direction concordance 6/6
Pooled clinical dose-response baseline	approved with limitations	dose bands described; pooled rows never treated as independent
Mechanistic proof of concept	conditionally approved	16 constructs over 12 independent sequence groups; trained: no
Sequence-only clinical classifier	BLOCKED	not run

The most distinctive content in this dataset is the matched-contrast set. Each contrast shares a nucleotide sequence and differs in exactly one named chemical factor — PS versus PO backbone, LNA wings added, or sequence alone — so the comparison isolates that

factor. Across 6/6 contrasts, this dataset and the project's independently curated scientific package agree on the *sign* of the effect; absolute grades agree in 1/6, and the disagreements are listed for scientific adjudication rather than silently reconciled.

What that agreement is, stated precisely — because an earlier version of this document overstated it. Both the scientific package's adjudicated score and this dataset's grade derive from the *same underlying publications*. Their agreement is therefore evidence that the extraction is **reproducible and consistently transcribed** by two independent curation efforts. It is **not** independent biological validation, and it is not external validation of any grade. Nothing in it replicates an experiment. An earlier draft described these as "two independently built datasets" agreeing, which is true of the curation and false of the evidence; the claim is corrected here.

The same correction applies to volume. Per-replicate numerical cells recovered from a source's own raw workbook are not independent observations: recovering them reconstructs that publication's reported measurements at finer grain. The staging manifest in [curation/research_staging/](#) states this on its face, and separates the count of cells from the count of constructs they describe.

What the contrasts do establish is narrower and still useful. The clearest is the phosphorothioate/phosphodiester pair of one 22-mer, where both curation efforts place the PS form at the top of the scale and its isosequential PO control at zero — the same nucleotide sequence, the backbone as the only variable. That is a *mechanistic* statement supported by a matched control, not a performance claim about a model. For a dataset whose purpose is to let others build models, a well-controlled mechanistic contrast with its raw measurements traceable is the more useful deliverable than a score.

Beyond this, the dataset is structured to support **dose–response modelling** from the banded clinical incidence grids, **cross-species comparison** on the 23-compound bridge (with exposures shown side by side and every row carrying an explicit statement of whether the comparison is qualified), and — the clinically decisive task — **separating the mild dose-dependent mode from the severe immune-mediated mode**, for which the `immunogenicity` readout category provides the labels. Each of those remains gated on scientific review, and none is claimed as achieved here.

What we do not claim. This is a feasibility demonstration, not a validated clinical predictor. Grade remains partly confounded with study type; the severe idiosyncratic mode is antibody-driven and is *not* expected to be predictable from design at all; and 32 % of rows have completed adversarial verification, so the remainder should be treated as curated-but-unverified. These limits are documented in the methodology rather than left for a reader to discover.