

Public Access & Dissemination Plan

OligoTox-Thrombocytopenia · NIH/NCATS OligoTox Open Data Challenge, Phase 2

This plan covers the **thrombocytopenia / platelet-toxicity** endpoint. It describes how the dataset is licensed, made publicly accessible and disseminated, and — as the Challenge requires — how the U.S. Government can allow interested parties to use it even if the submitting team does not itself continue to use or distribute it.

1. Scope — the solution covered by this plan

Artefact	Description
<code>data/oligos.csv</code>	254 oligonucleotides — identity and design predictors (20 columns), including sequence, per-residue modification map, phosphorothioate count and purity fields
<code>data/measurements.csv</code>	1,878 graded per-measurement platelet-toxicity records (24 columns) with per-row provenance and rights status
<code>data/measurements_human.csv · _animal.csv · bridge_human_animal.csv</code>	human (1,372) and animal (497) splits, plus the 22 compounds characterised on both sides
<code>data/oligotox_thrombo_merged.csv</code>	generated denormalised analysis view
<code>schema.md</code>	data dictionary, controlled vocabularies and the 0–3 grade rubric
<code>METHODOLOGY.md · SOURCES.md · README.md</code>	assembly and QC record, source registry with per-source rights, overview and findings
<code>scripts/</code>	the complete pipeline — assembly, verification, QC, reporting, splits, and the predictive-model demonstration
<code>curation/</code>	the full curation record: raw extractions, verification verdicts, source sweep

The dataset is a **curation of already-published data**. It contains **no human-subjects data, no personally identifiable information and no protected health information** — only aggregate experimental and clinical toxicology values and chemical design descriptors drawn from public literature, regulatory documents and patents. There are therefore **no privacy, consent or data-use-agreement constraints** on redistribution.

The curation record is included deliberately: it makes the published tables **reproducible from their inputs** rather than merely re-checkable against citations. It contains no third-party full texts — sources are referenced by identifier and exact locus, never redistributed — so it carries the same rights position as the tables themselves.

2. Licensing scheme

All material produced by the team — data tables, documentation, scripts and curation record — is released under the **Creative Commons Attribution 4.0 International (CC-BY 4.0)** licence: permissive, irrevocable, and permitting **anyone**, including the U.S. Government and any third party, to access, reuse, redistribute, modify and build upon the dataset in perpetuity subject only to attribution. A more permissive dedication such as CC0 1.0 can be substituted at NCATS's preference. These are **non-exclusive** licences for research and any other purpose.

Underlying third-party full texts are never redistributed. Rights are tracked **per row**, so a consumer can filter to exactly the records they may lawfully reuse:

Class	Meaning	Typical source	rows
public_domain	reproduce without restriction	USPTO patents; FDA and EMA documents	748
cc_by	raw values reproducible with attribution	CC-BY open-access articles, licence confirmed from the article's own licence field	203
summary_stat / derived_features_only	summary statistics or derived features under fair use	copyrighted journal content	927
verify	rights unresolved — settled before release	—	0

A `cc_by` classification is taken from the article's own licence field, never from the fact that it happens to be free to read. **No patents, trade secrets or restrictive intellectual property** are or will be claimed over the dataset, and there is no proprietary component whose withdrawal could remove public access. Because the licence is **irrevocable once granted**, third-party use does not depend on the team's continued participation.

3. Public access — hosting and persistence

- **Primary distribution:** a public code-hosting repository (`naviero1/Oligos`), openly accessible without registration or account.
- **Independent archival persistence and a citable identifier:** a versioned snapshot of each release will be deposited in a public long-term archive — **Zenodo** or an **NIH-designated repository** — and assigned a **DOI**. This guarantees the dataset remains findable and retrievable **independently of the team or any single hosting account**.
- **Non-proprietary formats:** UTF-8 CSV for data, Markdown for documentation, JSON for the curation record, Python for tooling. No software licence or paid dependency is needed to read or rebuild any of it.

4. Dissemination

Knowledge necessary to use the dataset is published with it, not held back: the data dictionary and schema, the methodology including grading conventions and QC, the source registry with per-source rights, the complete pipeline, and a worked predictive-model demonstration that shows how to load the data and train on it. Discoverability is supported by stable primary keys, descriptive metadata and keywords, the archival DOI, and links from the Challenge submission record.

FAIR principle	How this dataset meets it
Findable	Stable primary keys (<code>TOLG###</code> , <code>TMSR###</code>); archival DOI; descriptive metadata; public repository
Accessible	Open repository and archive; open formats; no gatekeeping, registration or request process
Interoperable	Controlled vocabularies; normalised two-table schema shared with a sister endpoint so the two join on oligo identity; standard formats
Reusable	CC-BY 4.0; per-row provenance (source identifier, reference, exact locus) so any value is re-verifiable; per-row rights so lawful reuse is decidable at the record level; committed inputs so the tables can be rebuilt

5. Continuity and U.S. Government use contingency

The Challenge requires a plan for how the U.S. Government can allow interested parties to use the solution if the team fails to use it and does not permit others to use it on reasonable terms. Three scenarios cover this, with the licensing scheme that applies to each:

Scenario	Condition	Licensing scheme employed
A — Normal operation	The team maintains the repository and archival deposit and disseminates updates.	CC-BY 4.0 to all third parties; no action required of any user beyond attribution.
B — Team ceases activity	The team stops maintaining or disseminating the dataset.	No action by the team is needed for continued public use. The dataset has already been released under an irrevocable CC-BY 4.0 licence and a copy is deposited in an independent public archive with a DOI . The Government and any interested party may continue to access, copy, host, modify and build upon it without further permission. The permissive licence and the independent archival copy are the mechanisms that accomplish this.
C — Explicit government grant	The Government wishes to authorise others directly.	The team grants the U.S. Government a non-exclusive, irrevocable, royalty-free right to use, reproduce, distribute, publicly display, host copies of, and authorise others to use the dataset and its documentation. This right survives any cessation of the team's activity.

In no scenario does third-party or Government use depend on the team's ongoing involvement, on a proprietary service, or on any revocable permission. The team agrees to abide by the terms of this plan as a condition of award, and understands that NIH intends to post or otherwise publicly display it.

6. Maintenance, versioning and honest status reporting

- **Versioned releases** with a changelog; each release archived under its own DOI so prior versions remain citable and retrievable.
- **Verification status is tracked in-repo and reported, not glossed.** Rows verified against their primary source carry a marker; the methodology states which blocks are verified and which remain outstanding. A release must not be described as fully verified while any block is outstanding.
- **Known limitations travel with the data**, including the purity-data limitation inherent to curating published results, the partial confounding of grade with study type, and the provisional status of grades pending subject-matter review.
- Corrections and additions are tracked through the repository's public history.

7. Stewardship and compliance

Maintainer: the submitting team, via the public repository, until and beyond archival deposit. **Fallback steward:** the independent public archive, which preserves the deposited snapshot and DOI irrespective of team status. **Challenge contact:** as listed in the Phase 2 submission record.

The plan is consistent with **NIH Scientific Data Sharing and Public Access policy** and with **FAIR**. An open, irrevocable licence ensures perpetual public usability; the absence of PII, PHI and human-subjects data removes any privacy or consent barrier to sharing; and per-row provenance and rights tracking make lawful reuse verifiable at the record level.