

Public Access & Dissemination Plan (PADP)

OligoTox-Kidney dataset — NIH/NCATS Oligonucleotide Toxicity (OligoTox) Open Data Challenge, Phase 2 (Data Generation).

This plan describes how the OligoTox-Kidney dataset is licensed, made publicly accessible, and disseminated, and — as required by the Challenge — **how the U.S. Government can allow interested parties to use the dataset even if the submitting team does not itself continue to use or distribute it**. It is a required component of the Phase 2 submission and is written to align with NIH Scientific Data Sharing / Public Access policy and the FAIR data principles.

1. Scope — the "solution" covered by this plan

The solution is the **curated, openly-releasable dataset** and its documentation:

Artifact	Description
data/oligos.csv	65 oligonucleotides — identity, sequence, purity/identity characterization + design predictors (20 columns)
data/measurements.csv	246 graded per-measurement nephrotoxicity records (25 columns), each carrying subject_class and renal_endpoints_measured
data/oligotox_kidney_merged.csv	denormalised analysis view, 246 × 46 (generated, not canonical)
data/human_animal_bridge.csv	15 oligos carrying paired human and animal evidence (generated)
LICENSE	CC BY 4.0 grant at the repository root, third-party source PDFs excluded
schema.md	Full data dictionary, controlled vocabularies, and the 0–3 grade rubric
METHODOLOGY.md	How the dataset was assembled (sources → extraction → grading → QC)
sources/SOURCES.md	Source registry (16 source IDs), redistribution status, acquisition state
README.md, PRESENTATION.md	Overview and findings deck

The dataset is a **curation of already-published data**. It contains **no human- subjects data, no personally identifiable information, and no protected health information** — only aggregate/experimental toxicology values and design descriptors drawn from public literature, regulatory labels, and patents. There are therefore **no privacy, consent, or data-use-agreement constraints** on redistribution.

2. Licensing scheme

- **Curated data tables and documentation** produced by the team are released under the **Creative Commons Attribution 4.0 International (CC-BY 4.0)** license — a **permissive, irrevocable** license that lets **anyone** (including the U.S. Government and any third party) access, reuse, redistribute, modify, and build upon the dataset in perpetuity, subject only to attribution. (*License choice is the team's proposal; a more permissive dedication such as CC0 1.0 can be substituted at NCATS's preference.*)
- **Underlying third-party full texts are never redistributed.** Journal articles are **referenced by DOI/PMID**, not copied. Values reproduced in the tables are limited to (a) **public-domain** sources — USPTO patents and U.S./EU regulatory labels — and (b) **summary statistics** used under fair use. Every row records its rights status in the redistribution column (public_domain or summary_stat).
- **No patents, trade secrets, or restrictive IP** are or will be claimed over the dataset. There is no proprietary component whose withdrawal could remove public access.

Because the license is **irrevocable once granted**, the ability of third parties to use the dataset does **not** depend on the team's continued participation.

3. Public access — hosting and persistence

- **Primary distribution:** a public code-hosting repository (GitHub — naviero1/Oligos), openly accessible without registration.
- **Independent archival persistence + citable identifier:** a versioned snapshot of each release will be deposited in a public, long-term archive (e.g., Zenodo or an NIH-designated repository) and assigned a DOI. This guarantees the dataset remains findable and retrievable **independently of the team or any single hosting account**.
- **Non-proprietary formats:** UTF-8 CSV for data and Markdown for documentation — openable with free, ubiquitous tooling, with no software dependency or license required to read.

4. Dissemination

- **Documentation for reuse:** README.md (overview), schema.md (data dictionary rubric), METHODOLOGY.md (assembly + QC), and the findings deck.
- **Discoverability:** descriptive metadata and keywords, a citable DOI, and links from the Challenge submission record.
- **FAIR alignment:**

Principle	How the dataset meets it
Findable	Stable primary keys (oligo_id, measurement_id); archival DOI; descriptive metadata
Accessible	Open repository + archive; open CSV/Markdown; no gatekeeping
Interoperable	Controlled vocabularies (common data elements), normalized two-table schema, standard formats
Reusable	CC-BY 4.0 license + per-row provenance (source_id, source_ref, source_table) so any value is re-verifiable

5. Continuity and U.S. Government use contingency (required)

The Challenge requires a plan for how the U.S. Government can allow interested parties to utilize the solution **if the team fails to utilize it and does not permit others to use it under reasonable terms**. The following scenarios cover this:

- **Scenario A — Normal operation.** The team maintains the repository and archival deposit and disseminates updates. Third parties use the dataset under CC-BY 4.0.
- **Scenario B — Team ceases activity or fails to disseminate.** No action by the team is needed for continued public use: the dataset has **already been released under an irrevocable CC-BY 4.0 license** and a **copy is deposited in an independent public archive with a DOI**. The U.S. Government and any interested party may continue to access, copy, host, and build upon the dataset **without further permission**. The permissive license and independent archival copy are the mechanisms that accomplish this plan.
- **Scenario C — Explicit government grant.** The team additionally grants the U.S. Government a **non-exclusive, irrevocable, royalty-free right** to use, reproduce, distribute, publicly display, host copies of, and **authorize others to use** the dataset and its documentation. This right survives any cessation of the team's activity and enables the Government to allow interested parties to utilize the dataset under reasonable terms.

In no scenario does third-party or Government use depend on the team's ongoing involvement, on a proprietary service, or on any revocable permission.

6. Maintenance and versioning

- **Versioned releases** with a changelog; each release archived under its own DOI so prior versions remain citable and retrievable.
 - **Provisional grades** are flagged in-row; a post-expert-review release will clear the `grade_provisional` flag and note the reviewer's disposition.
 - Corrections and additions are tracked through the repository's public history.
-

7. Stewardship

- **Maintainer:** the submitting team, via the public repository, until (and beyond) archival deposit.
 - **Fallback steward:** the independent public archive (Zenodo / NIH-designated repository) preserves the deposited snapshot and DOI irrespective of team status.
 - **Challenge contact:** as listed in the Phase 2 submission record.
-

8. Compliance summary

- Consistent with **NIH Scientific Data Sharing / Public Access policy** and **FAIR**.
- **Open, irrevocable license** (CC-BY 4.0) ensures perpetual public usability.
- **No PII/PHI / no human-subjects data** → no privacy or consent barriers to sharing.
- **Provenance and redistribution tracked per row** → lawful reuse is verifiable at the record level.